

Just text. Get things done.

An AI gig worker that lives in the messages you already send.

AI today talks. It doesn't do.

Every AI product wants you to install another app, learn another interface, chat in another tab. And after all that effort, the AI usually just... answers a question.

Real people don't want a chatbot. They want the bill paid, the photos organized, the party planned, the form submitted, the website built. They want it done in the channel they already live in: **SMS and email.**

**Text Gigler.
Or add it to any group
chat.
It picks up the goal and
runs.**

No app. No login. No prompt engineering. The interface is already in your pocket.

Every gig is its own AI agent.

A user texts a number or emails an address. Gigler spawns a dedicated agent for that conversation with its own memory, schedule, and toolset.

The agent picks the right tool: send a reply, draft a document, browse the web, run code in a sandbox, watch a video, schedule a follow-up, ask a human for approval.

Agent toolkit

- Reply via SMS / MMS / email
- Browse and fill out web forms
- Generate PDFs, sites, deliverables
- Run real code in a sandbox
- Read images, video, documents
- Schedule reminders & follow-ups
- Live View hand-off to a human

The team-parents group chat.

Saturday, 9am. A parent adds +1-555-GIGLER to the group thread:

- 9:02 AM** Mom: *"Hey Gigler, can you organize photos during the tournament and recap each game?"*
- 9:03 AM** **Gigler:** "On it. Drop photos and game clips here. I'll sort by player and recap when each game ends."
- 10-11 AM** Parents share photos and a livestream link. Gigler watches the video, identifies players via jersey numbers, tags photos automatically.
- 11:15 AM** **Gigler:** "Game 1 recap: 3-1 win. Top moments: Avery's assist (4:12), Maya's save (8:30). Photo album: gigler.ai/d/k7m2"
- 5:00 PM** **Gigler:** "Tournament wrap: 4 games, 32 photos sorted by 12 players, 4 recap videos, parent email sent. PDF album attached."

One channel. Everything.

Always on. Beach, office, or on the go. Text it, it's done.

FLAGSHIP CAPABILITY

Gig Programmer

Code, debug, and ship features. Anytime, anywhere. From a hammock if you want.

```
const app = express()
app.get('/api', handler)
app.listen(3000)
```

✓ Deployed successfully · preview at
gigler.ai/d/k7m2

Landing Pages

Full websites from a text description.

Event Planning

Book venues, send invites, coordinate.

Photo Collages

Organize and share event photos.

PDFs & Reports

Polished documents and summaries.

Restaurant Menus

Print-ready or shareable online.

Excel Sheets

Organize, analyze, share data.

Bills Dashboard

Auto-extracted utilities, due dates.

Workspace Workflows

Approvals, handoffs, automations.

Five things just clicked into place.

- LLMs are finally good enough for autonomous, multi-step execution. Not just chat.
- Vision + multimodal models can watch video, read receipts, classify uploads in real time.
- Headless browsers are programmable, observable, and human-in-the-loop ready.
- Sandboxes for code execution are cheap, safe, and per-session billable.
- SMS/MMS and email infrastructure is commodity. Reach is universal.

The first generation of consumer AI agents that actually do work, not just answer questions, is here.

Why we want to build on NVIDIA.

Today

- Google Gemini for reasoning, vision, function calling
- Twilio for SMS / MMS / Conversations
- Cloudflare Workers + Agents SDK + Durable Objects (in migration)
- Cloudflare R2 for media, Browser Run for web tasks
- Per-gig agent architecture with durable execution

What NVIDIA unlocks for us

- **NVIDIA NIM microservices** for cost-efficient, custom inference (vision tasks: jersey ID, receipt parsing, sports highlight detection)
- **NVIDIA Riva** for voice agents: phone-call gigs, voice memo ingestion, audio recap generation
- **TensorRT-LLM** for long-context, low-latency tool loops at scale
- **NeMo Guardrails** for safe autonomous execution in consumer hands
- **Inception ecosystem**: partners building on similar agent infrastructure

Live product. Active users. Aggressive roadmap.

7

ACTIVE TESTERS (ALPHA)

~\$1M

FUNDING RAISED

6 wks

TO CLOUDFLARE CUTOVER & V2
LAUNCH

- Live SMS & email gig completion at gigler.ai today.
- Sister company Carmen.ai validates the underlying multi-agent infrastructure.
- Roadmap: per-gig durable agents (Q2), code sandboxes (Q3), browser agents (Q3), voice agents on NVIDIA Riva (Q4).

Let's build the SMS-native AI agent on NVIDIA.

We're applying to NVIDIA Inception for compute credits, technical access to NIM and Riva, and a partnership path to ship voice + vision agents at consumer scale. Gigler is the everyday-user interface NVIDIA agents have been waiting for.

gigler.ai

support@carmenai.co